

Practical 1

Python Code:

```
import pandas as pd
import numpy as np
df=pd.read_csv("Iris.csv")
print(df.head())
df.head()
df.tail()
df.describe()
df.info()
df.shape
df.dtypes
df.columns
missing_values=df.isnull().sum()
print("\nMissing values in csv data:")
print("Missing_values")
df = df.drop(columns=['Species'])
fill= df.fillna(df.mean())
print("\nFilled Missing Values with Mean (CSV Data):")
print(df.head())
z_scores = np.abs((fill - fill.mean()) / fill.std())
z_scores
do = fill[(z_scores < 3).all(axis=1)]
print("\nData After Removing Outliers (CSV Data):")
print(do.head())
threshold_value = 3
filter = do[do['SepalLengthCm'] > threshold_value]
print(f"\nFiltered Data (Rows with column_name > {threshold_value}):")
print(filter)
df_sorted = filter.sort_values(by='SepalWidthCm', ascending=False)
print(df_sorted.head())
df_grouped = df_sorted.groupby('SepalLengthCm').agg({
    'PetalLengthCm': 'mean',
    'PetalWidthCm': 'sum'
}).reset_index()
print("\nGrouped Data (Mean and Sum for Each Group):")
print(df_grouped.head())
```

Iris.csv:

Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
1	5.1	3.5	1.4	0.2	Iris-setosa
2	4.9	3.0	1.4	0.2	Iris-setosa
3	4.7	3.2	1.3	0.2	Iris-setosa
4	4.6	3.1	1.5	0.2	Iris-setosa
5	5.0	3.6	1.4	0.2	Iris-setosa
6	5.4	3.9	1.7	0.4	Iris-setosa
7	4.6	3.4	1.4	0.3	Iris-setosa

```

8,5.0,3.4,1.5,0.2,Iris-setosa
9,4.4,2.9,1.4,0.2,Iris-setosa
10,4.9,3.1,1.5,0.1,Iris-setosa
11,5.4,3.0,4.5,1.5,Iris-versicolor
12,6.0,3.4,4.5,1.6,Iris-versicolor
13,6.7,3.1,4.7,1.5,Iris-versicolor
14,6.3,2.3,4.4,1.3,Iris-versicolor
15,5.6,3.0,4.1,1.3,Iris-versicolor
16,5.5,2.5,4.0,1.3,Iris-versicolor
17,5.5,2.6,4.4,1.2,Iris-versicolor
18,6.1,3.0,4.6,1.4,Iris-versicolor
19,5.8,2.6,4.0,1.2,Iris-versicolor
20,5.0,2.3,3.3,1.0,Iris-versicolor
21,6.3,3.3,6.0,2.5,Iris-virginica
22,5.8,2.7,5.1,1.9,Iris-virginica
23,7.1,3.0,5.9,2.1,Iris-virginica
24,6.3,2.9,5.6,1.8,Iris-virginica
25,6.5,3.0,5.2,2.0,Iris-virginica
26,7.6,3.0,6.6,2.1,Iris-virginica
27,4.9,2.5,4.5,1.7,Iris-virginica
28,7.3,2.9,6.3,1.8,Iris-virginica
29,6.7,2.5,5.8,1.8,Iris-virginica
30,7.2,3.6,6.1,2.5,Iris-virginica

```

OUTPUT:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

<class 'pandas.core.frame.DataFrame'>

Rангінг: 30 entries, 0 to 29

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	Id	30 non-null	int64
1	SepalLengthCm	30 non-null	float64
2	SepalWidthCm	30 non-null	float64
3	PetalLengthCm	30 non-null	float64
4	PetalWidthCm	30 non-null	float64
5	Species	30 non-null	object

dtypes: float64(4), int64(1), object(1)

memory usage: 1.5+ KB

Missing values in csv data:

Missing_values

Filled Missing Values with Mean (CSV Data):

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	1	5.1	3.5	1.4	0.2
1	2	4.9	3.0	1.4	0.2
2	3	4.7	3.2	1.3	0.2
3	4	4.6	3.1	1.5	0.2
4	5	5.0	3.6	1.4	0.2

Data After Removing Outliers (CSV Data):

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	1	5.1	3.5	1.4	0.2
1	2	4.9	3.0	1.4	0.2
2	3	4.7	3.2	1.3	0.2
3	4	4.6	3.1	1.5	0.2
4	5	5.0	3.6	1.4	0.2

Filtered Data (Rows with column_name > 3):

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	1	5.1	3.5	1.4	0.2
1	2	4.9	3.0	1.4	0.2
2	3	4.7	3.2	1.3	0.2
3	4	4.6	3.1	1.5	0.2
4	5	5.0	3.6	1.4	0.2
5	6	5.4	3.9	1.7	0.4
6	7	4.6	3.4	1.4	0.3
7	8	5.0	3.4	1.5	0.2
8	9	4.4	2.9	1.4	0.2
9	10	4.9	3.1	1.5	0.1
10	11	5.4	3.0	4.5	1.5
11	12	6.0	3.4	4.5	1.6
12	13	6.7	3.1	4.7	1.5
13	14	6.3	2.3	4.4	1.3
14	15	5.6	3.0	4.1	1.3
15	16	5.5	2.5	4.0	1.3
16	17	5.5	2.6	4.4	1.2
17	18	6.1	3.0	4.6	1.4
18	19	5.8	2.6	4.0	1.2
19	20	5.0	2.3	3.3	1.0
20	21	6.3	3.3	6.0	2.5
21	22	5.8	2.7	5.1	1.9
22	23	7.1	3.0	5.9	2.1
23	24	6.3	2.9	5.6	1.8
24	25	6.5	3.0	5.2	2.0
25	26	7.6	3.0	6.6	2.1
26	27	4.9	2.5	4.5	1.7
27	28	7.3	2.9	6.3	1.8
28	29	6.7	2.5	5.8	1.8
29	30	7.2	3.6	6.1	2.5

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
5	6	5.4	3.9	1.7	0.4
4	5	5.0	3.6	1.4	0.2
29	30	7.2	3.6	6.1	2.5
0	1	5.1	3.5	1.4	0.2
7	8	5.0	3.4	1.5	0.2

Grouped Data (Mean and Sum for Each Group):

	SepalLengthCm	PetalLengthCm	PetalWidthCm
0	4.4	1.400000	0.2
1	4.6	1.450000	0.5
2	4.7	1.300000	0.2
3	4.9	2.466667	2.0
4	5.0	2.066667	1.4

Practical 2

Python Code:

```
#part1:feature scaling
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder
df = pd.read_csv("wine.csv")
print("\nFull CSV Loaded:")
print(df.head())
df1 = pd.read_csv("wine.csv", usecols=[0, 1, 2], skiprows=1)
df1.columns = ['classlabel', 'Alcohol', 'Malic Acid']
print("\nOriginal DataFrame (First 5 Rows):")
print(df1.head())
scaling = MinMaxScaler()
scaled_value = scaling.fit_transform(df1[['Alcohol', 'Malic Acid']])
df1[['Alcohol', 'Malic Acid']] = scaled_value
print("\nDataFrame After Min-Max Scaling:")
print(df1.head())
scaling = StandardScaler()
scaled_standardvalue = scaling.fit_transform(df1[['Alcohol', 'Malic Acid']])
df1[['Alcohol', 'Malic Acid']] = scaled_standardvalue
print("\nDataFrame After Standard Scaling:")
print(df1.head())
plt.scatter(df1['Alcohol'], df1['Malic Acid'])
plt.title("Standard Scaled Alcohol vs Malic Acid")
plt.xlabel("Alcohol")
plt.ylabel("Malic Acid")
plt.show()

#part2:feature dummification
iris=pd.read_csv("Iris.csv")
iris
le=LabelEncoder()
iris['code']=le.fit_transform(iris.Species)
iris
wine.csv
```

classlabel	Alcohol	Malic Acid
1,14.23,1.71		
1,13.20,1.78		
1,13.16,2.36		
1,14.37,1.95		
1,13.24,2.59		
2,12.37,1.13		
2,12.33,1.10		
2,13.30,1.40		
2,12.71,1.45		
2,12.00,1.38		
3,13.75,1.75		
3,14.10,2.20		
3,14.00,1.90		
3,13.90,2.30		
3,13.50,1.60		

Output

Full CSV Loaded:

	classlabel	Alcohol	Malic Acid
0	1	14.23	1.71
1	1	13.20	1.78
2	1	13.16	2.36
3	1	14.37	1.95
4	1	13.24	2.59

Original DataFrame (First 5 Rows):


	classlabel	Alcohol	Malic Acid
0	1	13.20	1.78
1	1	13.16	2.36
2	1	14.37	1.95
3	1	13.24	2.59
4	2	12.37	1.13

Dataframe After Min-Max Scaling:

	classlabel	Alcohol	Malic Acid
0	1	0.506329	0.456376
1	1	0.489451	0.845638
2	1	1.000000	0.570470
3	1	0.523207	1.000000
4	2	0.156118	0.020134

Dataframe After Standard Scaling:

	classlabel	Alcohol	Malic Acid
0	1	-0.116248	0.004780
1	1	-0.173858	1.298461
2	1	1.568833	0.383962
3	1	-0.058638	1.811472
4	2	-1.311647	-1.445035



Practical 3

Python Code:

```
#part1
import pandas as pd
import numpy as np
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
df = pd.read_csv("FCH.csv")
print("\nCSV Data Loaded:")
print(df.head())
print(df.tail())
print("Shape:", df.shape)
print("Size:", df.size)
print(df.describe())
print(df.info())
print("\nData Types:\n", df.dtypes)
housing = fetch_california_housing()
housing_df = pd.DataFrame(housing.data, columns=housing.feature_names)
housing_df["PRICE"] = housing.target
print("\nCalifornia Housing Dataset:")
print(housing_df.head())
#part2
X = housing_df[['AveRooms']]
y = housing_df[['PRICE']]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = LinearRegression()
model.fit(X_train, y_train)
mse = mean_squared_error(y_test, model.predict(X_test))
r2 = r2_score(y_test, model.predict(X_test))
print("\n===== SIMPLE LINEAR REGRESSION =====")
print("Mean Squared Error:", mse)
print("R-squared Value:", r2)
print("Intercept:", model.intercept_)
print("Coefficient:", model.coef_)

#part2:
X = housing_df.drop("PRICE", axis=1)
y = housing_df["PRICE"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```

print("\n===== MULTIPLE LINEAR REGRESSION =====")
print("Mean Squared Error:", mse)
print("R-squared Value:", r2)
print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)

```

FCH.csv:

MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	PRICE
8.3252	41	6.9841	1.0238	322	2.5556	37.88	-122.23	4.526
8.3014	21	6.2381	0.9719	2401	2.1098	37.86	-122.22	3.585
7.2574	52	8.2881	1.0734	496	2.8023	37.85	-122.24	3.521
5.6431	52	5.8174	1.0734	558	2.5479	37.84	-122.25	3.413
3.8462	52	6.2810	1.0811	565	2.1815	37.83	-122.26	3.422
4.0368	42	6.0836	1.0542	413	2.2426	37.85	-122.28	2.697
3.6591	52	6.0080	1.0234	1094	2.1270	37.85	-122.29	1.854
3.1200	52	5.0000	1.0200	850	2.1100	37.85	-122.30	2.234
2.8450	20	4.5000	0.9800	700	2.2000	37.80	-122.32	1.965
3.5000	25	6.2000	1.0500	900	2.3000	37.78	-122.33	2.540

OUTPUT:

CSV Data Loaded:

	MedInc	HouseAge	AveRooms	AveBedrms	...	AveOccup	Latitude	Longitude	PRICE
0	8.3252	41	6.9841	1.0238	...	2.5556	37.88	-122.23	4.526
1	8.3014	21	6.2381	0.9719	...	2.1098	37.86	-122.22	3.585
2	7.2574	52	8.2881	1.0734	...	2.8023	37.85	-122.24	3.521
3	5.6431	52	5.8174	1.0734	...	2.5479	37.84	-122.25	3.413
4	3.8462	52	6.2810	1.0811	...	2.1815	37.83	-122.26	3.422

[5 rows x 9 columns]

	MedInc	HouseAge	AveRooms	AveBedrms	...	AveOccup	Latitude	Longitude	PRICE
5	4.0368	42	6.0836	1.0542	...	2.2426	37.85	-122.28	2.697
6	3.6591	52	6.0080	1.0234	...	2.1270	37.85	-122.29	1.854
7	3.1200	52	5.0000	1.0200	...	2.1100	37.85	-122.30	2.234
8	2.8450	20	4.5000	0.9800	...	2.2000	37.80	-122.32	1.965
9	3.5000	25	6.2000	1.0500	...	2.3000	37.78	-122.33	2.540

[5 rows x 9 columns]

Shape: (10, 9)

Size: 90

	MedInc	HouseAge	AveRooms	...	Latitude	Longitude	PRICE
count	10.000000	10.000000	10.000000	...	10.000000	10.000000	10.000000
mean	5.053420	40.900000	6.140030	...	37.839000	-122.272000	2.975700
std	2.158254	13.755403	1.025986	...	0.029231	0.037947	0.853737
min	2.845000	20.000000	4.500000	...	37.780000	-122.330000	1.854000
25%	3.539775	29.000000	5.865050	...	37.832500	-122.297500	2.310500
50%	3.941500	47.000000	6.141800	...	37.850000	-122.270000	3.055000
75%	6.853825	52.000000	6.270275	...	37.850000	-122.242500	3.496250
max	8.325200	52.000000	8.288100	...	37.880000	-122.220000	4.526000

[8 rows x 9 columns]

<class 'pandas.core.frame.DataFrame'>

```
RangeIndex: 10 entries, 0 to 9  
Data columns (total 9 columns):  
 # Column Non-Null Count Dtype  
---  
 0 MedInc    10 non-null   float64  
 1 HouseAge   10 non-null   int64  
 2 AveRooms   10 non-null   float64  
 3 AveBedrms  10 non-null   float64  
 4 Population 10 non-null   int64  
 5 AveOccup   10 non-null   float64  
 6 Latitude   10 non-null   float64  
 7 Longitude  10 non-null   float64  
 8 PRICE      10 non-null   float64  
dtypes: float64(7), int64(2)  
memory usage: 852.0 bytes
```

```
None
```

```
Data Types:
```

```
MedInc    float64  
HouseAge   int64  
AveRooms  float64  
AveBedrms float64  
Population int64  
AveOccup   float64  
Latitude   float64  
Longitude  float64  
PRICE     float64  
dtype: object
```

```
California Housing Dataset:
```

```
MedInc HouseAge AveRooms AveBedrms ... AveOccup Latitude Longitude PRICE  
0 8.3252  41.0 6.984127 1.023810 ... 2.555556  37.88 -122.23 4.526  
1 8.3014  21.0 6.238137 0.971880 ... 2.109842  37.86 -122.22 3.585  
2 7.2574  52.0 8.288136 1.073446 ... 2.802260  37.85 -122.24 3.521  
3 5.6431  52.0 5.817352 1.073059 ... 2.547945  37.85 -122.25 3.413  
4 3.8462  52.0 6.281853 1.081081 ... 2.181467  37.85 -122.25 3.422
```

```
[5 rows x 9 columns]
```

```
===== SIMPLE LINEAR REGRESSION =====
```

```
Mean Squared Error: 1.2923314440807299
```

```
R-squared Value: 0.013795337532284901
```

```
Intercept: [1.65476227]
```

```
Coefficient: [[0.07675559]]
```

```
===== MULTIPLE LINEAR REGRESSION =====
```

```
Mean Squared Error: 0.555891598695244
```

```
R-squared Value: 0.5757877060324511
```

```
Intercept: -37.023277706064064
```

```
Coefficients: [ 4.48674910e-01 9.72425752e-03 -1.23323343e-01 7.83144907e-01
```

```
-2.02962058e-06 -3.52631849e-03 -4.19792487e-01 -4.33708065e-01]
```

Practical 4

Python Code:

```
#part1
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
df = pd.read_csv("Iris.csv")
print("\nOriginal Dataset:")
print(df.head())
df = df[df['Species'] != 'Iris-virginica']
print("\nDataset After Removing Class 2:")
print(df.head())
X = df.drop('Species', axis=1)
y = df['Species']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
y_pred_logistic = logistic_model.predict(X_test)
print("\n===== Logistic Regression Results =====")
print("Accuracy:", accuracy_score(y_test, y_pred_logistic))
print("\nClassification Report:\n", classification_report(y_test, y_pred_logistic))
#part2
from sklearn.tree import DecisionTreeClassifier
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
y_pred_tree = model.predict(X_test)
print("\n===== Decision Tree Classifier Results =====")
print("Accuracy:", accuracy_score(y_test, y_pred_tree))
print("\nClassification Report:\n", classification_report(y_test, y_pred_tree))
```

OUTPUT:

Original Dataset:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

Dataset After Removing Class 2:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa

```

2 3      4.7      3.2      1.3      0.2 Iris-setosa
3 4      4.6      3.1      1.5      0.2 Iris-setosa
4 5      5.0      3.6      1.4      0.2 Iris-setosa

```

===== Logistic Regression Results =====

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	2
Iris-versicolor	1.00	1.00	1.00	2
accuracy			1.00	4
macro avg	1.00	1.00	1.00	4
weighted avg	1.00	1.00	1.00	4

===== Decision Tree Classifier Results =====

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	2
Iris-versicolor	1.00	1.00	1.00	2
accuracy			1.00	4
macro avg	1.00	1.00	1.00	4
weighted avg	1.00	1.00	1.00	4

Practical 5

Python Code:

```

import warnings
warnings.filterwarnings("ignore")
import os
os.environ["LOKY_MAX_CPU_COUNT"] = "1"
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.datasets import make_blobs
X, _ = make_blobs(n_samples=300, centers=4, random_state=42)
inertia = []
K_range = range(1, 11)
for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X)
    inertia.append(kmeans.inertia_)
plt.plot(K_range, inertia, marker='o', linestyle='--')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow Method for Optimal k')
plt.show()
kmeans = KMeans(n_clusters=4, random_state=42, n_init=10)
y_kmeans = kmeans.fit_predict(X)
plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, cmap='viridis', edgecolors='k')
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],


```

```

        s=200, marker='X', label='Centroids')
plt.title('K-Means Clustering Result')
plt.legend()
plt.show()

```

OUTPUT:



Practical 6

Python Code:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
data = load_iris()
X = data.data
y = data.target
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
pca = PCA()
X_pca = pca.fit_transform(X_scaled)
explained_variance = pca.explained_variance_ratio_
cumulative_variance = np.cumsum(explained_variance)
plt.figure(figsize=(6, 4))
plt.plot(range(1, len(explained_variance) + 1), cumulative_variance,
         marker='o', linestyle='--')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Explained Variance vs Number of Components')
plt.grid(True)
plt.show()
pca_2d = PCA(n_components=2)
X_pca_2d = pca_2d.fit_transform(X_scaled)
plt.figure(figsize=(6, 4))



```

```

plt.scatter(X_pca_2d[:, 0], X_pca_2d[:, 1], c=y, cmap='viridis', edgecolors='k')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('PCA: 2D Projection of Iris Dataset')
plt.colorbar(label='Class Label')
plt.grid(True)
plt.show()

```

OUTPUT:



Practical 8

Python Code:

```

import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
np.random.seed(42)

sample1 = np.random.normal(loc=10, scale=2, size=30)
sample2 = np.random.normal(loc=12, scale=2, size=30)
t_statistic, p_value = stats.ttest_ind(sample1, sample2)
alpha = 0.05
print("Results of Two-Sample t-test:")
print("T-statistic:", t_statistic)
print("P-value:", p_value)
print("Degrees of Freedom:", len(sample1) + len(sample2) - 2)
plt.figure(figsize=(10, 5))
plt.hist(sample1, alpha=0.5, label='Sample 1', color='blue')
plt.hist(sample2, alpha=0.5, label='Sample 2', color='orange')
plt.axvline(np.mean(sample1), color='blue', linestyle='dashed', linewidth=2)
plt.axvline(np.mean(sample2), color='orange', linestyle='dashed', linewidth=2)
plt.title('Distributions of Sample 1 and Sample 2')
plt.xlabel('Values')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True)
t_critical = stats.t.ppf(1 - alpha/2, len(sample1) + len(sample2) - 2)

```

```

sample_min, sample_max = min(sample1.min(), sample2.min()), max(sample1.max(), sample2.max())
x_vals = np.linspace(sample_min, sample_max, 100)
critical_region = np.abs(x_vals - np.mean(sample1)) > t_critical
plt.fill_between(x_vals, 0, 5, where=critical_region, alpha=0.3, label='Critical Region')
plt.text(11.5, 5, f"t = {t_statistic:.2f}", ha='center', color='black', backgroundcolor='white')
plt.show()

if p_value < alpha:
    print("\nConclusion: There is significant evidence to reject the null hypothesis.")
    print("Interpretation: The mean of Sample 1 is significantly higher than that of Sample 2.")
else:
    print("\nConclusion: Fail to reject the null hypothesis.")
    print("Interpretation: There is not enough evidence to claim a significant difference between the means.")

```

```

# chi-test
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import warnings
from scipy import stats
warnings.filterwarnings('ignore')
df = sb.load_dataset('mpg')
print(df)
print(df['horsepower'].describe())
print(df['model_year'].describe())
# Binning horsepower
bins = [0, 75, 150, 240]
df['horsepower_new'] = pd.cut(df['horsepower'], bins=bins, labels=['l', 'm', 'h'])
print(df['horsepower_new'])
ybins = [69, 72, 74, 84]
label = ['1', '2', '3']
df['modelyear_new'] = pd.cut(df['model_year'], bins=ybins, labels=label)
newyear = df['modelyear_new']
print(newyear)
df_chi = pd.crosstab(df['horsepower_new'], df['modelyear_new'])
print(df_chi)
print(stats.chi2_contingency(df_chi))

```

OUTPUT:

Results of Two-Sample t-test:
T-statistic: -4.512913234547555
P-value: 3.176506547470154e-05
Degrees of Freedom: 58

Conclusion: There is significant evidence to reject the null hypothesis.

Interpretation: The mean of Sample 1 is significantly higher than that of Sample 2.

```
mpg cylinders ... origin name
0 18.0     8 ... usa  chevrolet chevelle malibu
1 15.0     8 ... usa  buick skylark 320
2 18.0     8 ... usa  plymouth satellite
3 16.0     8 ... usa  amc rebel sst
4 17.0     8 ... usa  ford torino
...
393 27.0     4 ... usa  ford mustang gl
394 44.0     4 ... europe  vw pickup
395 32.0     4 ... usa  dodge rampage
396 28.0     4 ... usa  ford ranger
397 31.0     4 ... usa  chevy s-10
```

[398 rows x 9 columns]

count 392.000000

mean 104.469388

std 38.491160

min 46.000000

25% 75.000000

50% 93.500000

75% 126.000000

max 230.000000

Name: horsepower, dtype: float64

count 398.000000

mean 76.010050

std 3.697627

min 70.000000

25% 73.000000

50% 76.000000

75% 79.000000

max 82.000000

Name: model_year, dtype: float64

0 m

1 h

2 m

3 m

4 m

..

393 m

394 l

395 m

396 m

397 m

Name: horsepower_new, Length: 398, dtype: category

Categories (3, object): ['l' < 'm' < 'h']

0 1

1 1

2 1

```

3    1
4    1
..
393   3
394   3
395   3
396   3
397   3
Name: modelyear_new, Length: 398, dtype: category
Categories (3, object): ['1' < '2' < '3']
modelyear_new  1  2  3
horsepower_new
l      9 14 76
m     49 41 158
h     26 11  8
Chi2ContingencyResult(statistic=np.float64(54.95485392447537),
pvalue=np.float64(3.320518009555984e-11), dof=4, expected_freq=array([[ 21.21428571,  16.66836735,
61.11734694],
[ 53.14285714, 41.75510204, 153.10204082],
[ 9.64285714, 7.57653061, 27.78061224]]))

```

Figure 1

